

SELECTING THE BEST MODEL TO PREDICT HOUSING PURCHASING AMOUNT USING REGRESSION

BUDI SALEH, Created on 24 Sep 20

By using the boston dataset from Kagle and we would like to develop a model to predict the total price housing amount that customers are willing to pay given the following attributes:

- No of bedrooms
- square foot (living room,lot, above, and basement)
- City
- View
- Condition
- Year built and renovated
- Bathrooms

The model should predict:

- Price of House

This exercise to find the best regression model based on the

- Mean square error
- Mean absolute error
- r2 score

Linear Model are using as follow :

1. Linear regression as baseline
2. Ridge
3. Lasso
4. KNN
5. Decission Tree
6. Random Forrest
7. Support Vector

Import the library

```
In [1]: import matplotlib as mp
import sklearn
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import cross_val_score, train_test_split, GridSearchCV
from sklearn.metrics import mean_squared_error as mse
from sklearn.metrics import r2_score as rs
from sklearn.metrics import mean_absolute_error as mae
import category_encoders as ce
from math import sqrt
from sklearn.svm import SVR # Support Vector Regressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeRegressor
from sklearn.svm import LinearSVR
# to perform hyperparameter tuning
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
```

```
In [2]: %matplotlib inline
```

Using data download from Kagle for Boston Housing dataset - Read using panda


```
In [3]: df = pd.read_csv('data.csv')
```

In [4]: df

Out[4]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
0	2014-05-02 00:00:00	3.130000e+05	3.0	1.50	1340	7912	1.5	0	
1	2014-05-02 00:00:00	2.384000e+06	5.0	2.50	3650	9050	2.0	0	
2	2014-05-02 00:00:00	3.420000e+05	3.0	2.00	1930	11947	1.0	0	
3	2014-05-02 00:00:00	4.200000e+05	3.0	2.25	2000	8030	1.0	0	
4	2014-05-02 00:00:00	5.500000e+05	4.0	2.50	1940	10500	1.0	0	
...
4595	2014-07-09 00:00:00	3.081667e+05	3.0	1.75	1510	6360	1.0	0	
4596	2014-07-09 00:00:00	5.343333e+05	3.0	2.50	1460	7573	2.0	0	
4597	2014-07-09 00:00:00	4.169042e+05	3.0	2.50	3010	7014	2.0	0	
4598	2014-07-10 00:00:00	2.034000e+05	4.0	2.00	2090	6630	1.0	0	
4599	2014-07-10 00:00:00	2.206000e+05	3.0	2.50	1490	8102	2.0	0	

4600 rows × 18 columns



Exploratory Data Analysis

In [5]: df.describe()

Out[5]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
count	4.600000e+03	4600.000000	4600.000000	4600.000000	4.600000e+03	4600.000000	4600.000000
mean	5.519630e+05	3.400870	2.160815	2139.346957	1.485252e+04	1.512065	0.000000
std	5.638347e+05	0.908848	0.783781	963.206916	3.588444e+04	0.538288	0.000000
min	0.000000e+00	0.000000	0.000000	370.000000	6.380000e+02	1.000000	0.000000
25%	3.228750e+05	3.000000	1.750000	1460.000000	5.000750e+03	1.000000	0.000000
50%	4.609435e+05	3.000000	2.250000	1980.000000	7.683000e+03	1.500000	0.000000
75%	6.549625e+05	4.000000	2.500000	2620.000000	1.100125e+04	2.000000	0.000000
max	2.659000e+07	9.000000	8.000000	13540.000000	1.074218e+06	3.500000	1.000000

In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  4600 non-null  object
1   price                 4600 non-null  float64
2   bedrooms              4600 non-null  float64
3   bathrooms             4600 non-null  float64
4   sqft_living           4600 non-null  int64
5   sqft_lot              4600 non-null  int64
6   floors                4600 non-null  float64
7   waterfront            4600 non-null  int64
8   view                  4600 non-null  int64
9   condition             4600 non-null  int64
10  sqft_above            4600 non-null  int64
11  sqft_basement         4600 non-null  int64
12  yr_built              4600 non-null  int64
13  yr_renovated          4600 non-null  int64
14  street                4600 non-null  object
15  city                  4600 non-null  object
16  statezip              4600 non-null  object
17  country               4600 non-null  object
dtypes: float64(4), int64(9), object(5)
memory usage: 647.0+ KB
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: date           0
price           0
bedrooms        0
bathrooms       0
sqft_living     0
sqft_lot        0
floors          0
waterfront      0
view            0
condition       0
sqft_above      0
sqft_basement   0
yr_built        0
yr_renovated    0
street          0
city            0
statezip        0
country         0
dtype: int64
```

Drop unnecessary column

```
In [8]: df2 = df.drop(['date', 'street', 'statezip', 'country'], axis =1)
```

```
In [9]: df2
```

```
Out[9]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	conditi
0	3.130000e+05	3.0	1.50	1340	7912	1.5	0	0	
1	2.384000e+06	5.0	2.50	3650	9050	2.0	0	4	
2	3.420000e+05	3.0	2.00	1930	11947	1.0	0	0	
3	4.200000e+05	3.0	2.25	2000	8030	1.0	0	0	
4	5.500000e+05	4.0	2.50	1940	10500	1.0	0	0	
...	
4595	3.081667e+05	3.0	1.75	1510	6360	1.0	0	0	
4596	5.343333e+05	3.0	2.50	1460	7573	2.0	0	0	
4597	4.169042e+05	3.0	2.50	3010	7014	2.0	0	0	
4598	2.034000e+05	4.0	2.00	2090	6630	1.0	0	0	
4599	2.206000e+05	3.0	2.50	1490	8102	2.0	0	0	

4600 rows × 14 columns

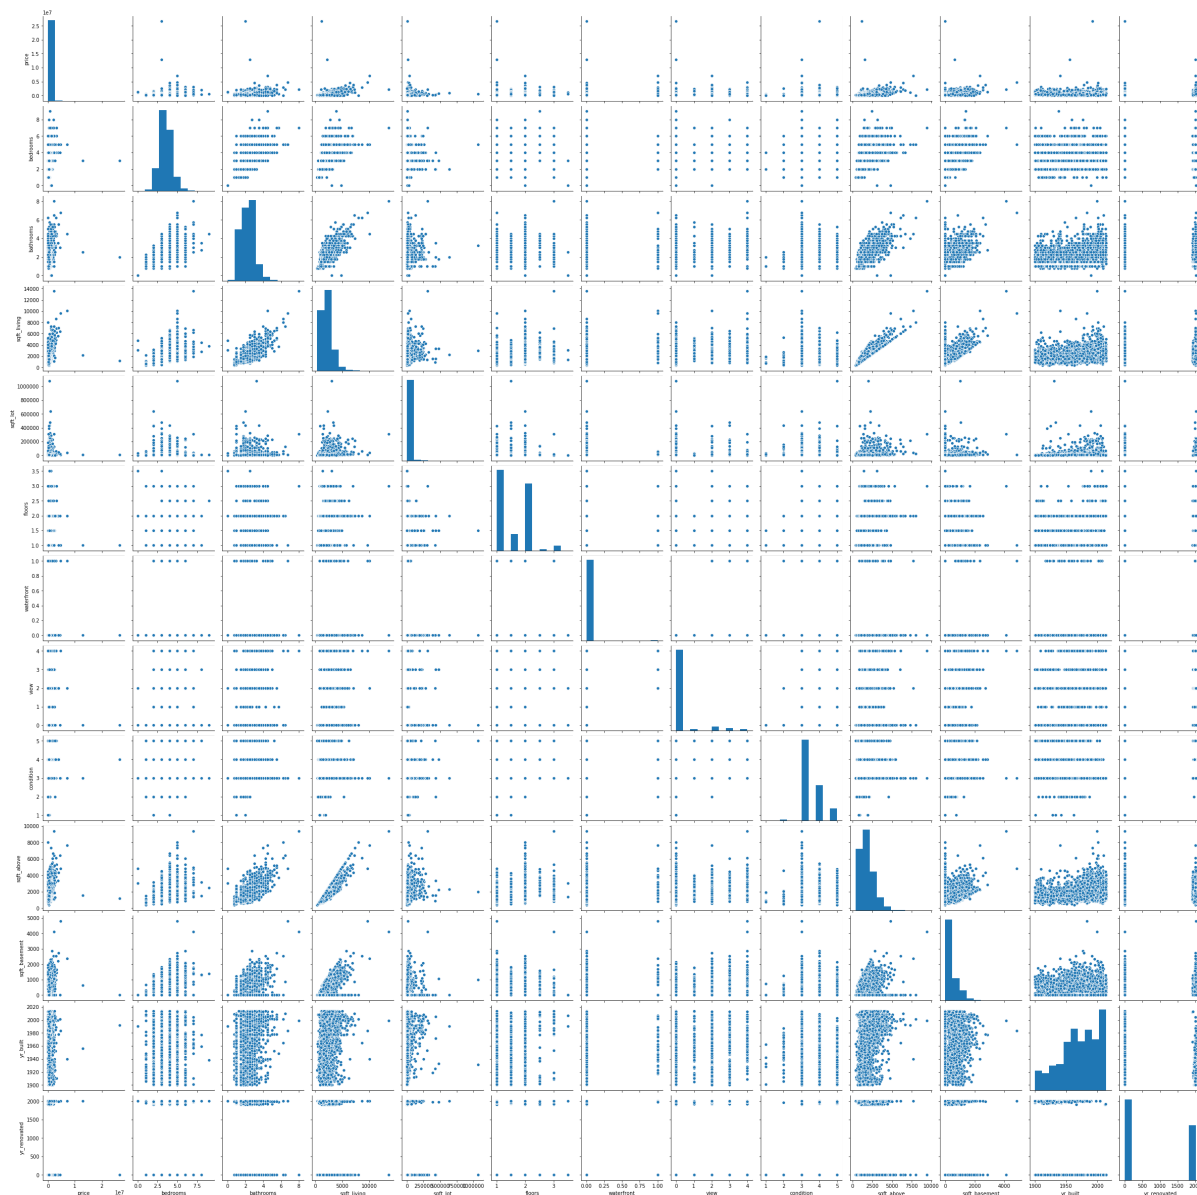


Visual the data

Using pairplot to see the correlation

```
In [10]: sns.pairplot(df2)
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x5493dc8>
```



From the graph, it can be seen that there are relations with regression

Create a correlation heatmap

```
In [11]: corr = df2.corr()
```

```
In [12]: ax = sns.heatmap(
    corr,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(20, 220, n=200),
    square=True
)
ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation=45,
    horizontalalignment='right'
);
```



Correlation table

```
In [13]: cmap = cmap=sns.diverging_palette(5, 250, as_cmap=True)

def magnify():
    return [dict(selector="th",
                  props=[("font-size", "7pt")]),
            dict(selector="td",
                  props=[('padding', "0em 0em")]),
            dict(selector="th:hover",
                  props=[("font-size", "12pt")]),
            dict(selector="tr:hover td:hover",
                  props=[('max-width', '200px'),
                          ('font-size', '12pt')])]

corr.style.background_gradient(cmap, axis=1)\
    .set_properties(**{'max-width': '80px', 'font-size': '10pt'})\
    .set_caption("Correlation Table")\
    .set_precision(2)\
    .set_table_styles(magnify())
```

Out[13]:

Correlation Table

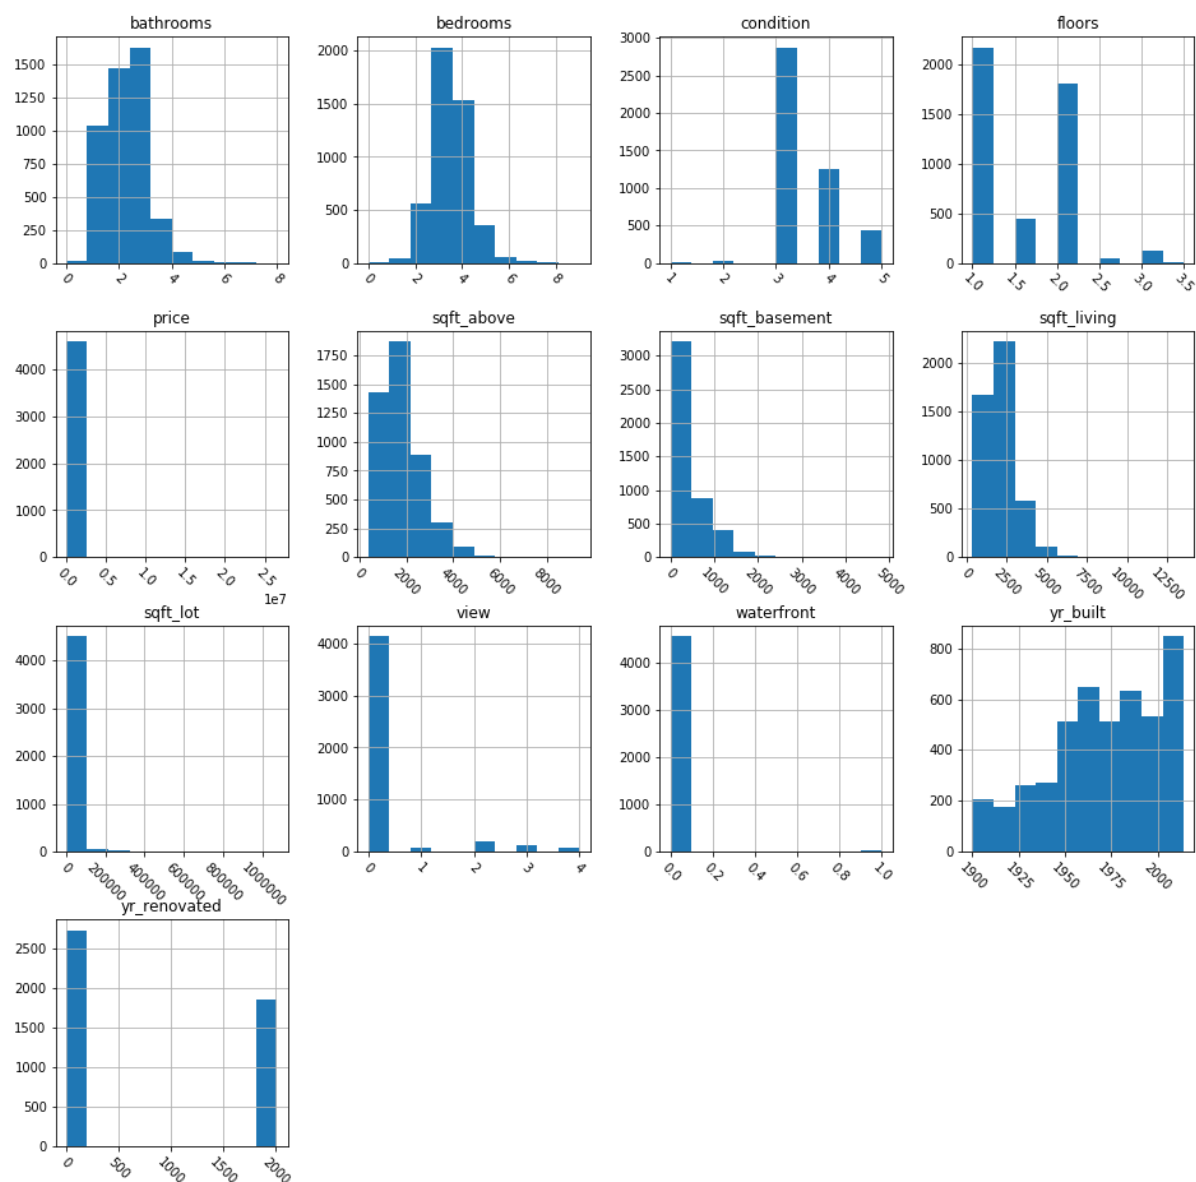
	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_base
price	1.00	0.20	0.33	0.43	0.05	0.15	0.14	0.23	0.03	0.37	
bedrooms	0.20	1.00	0.55	0.59	0.07	0.18	-0.00	0.11	0.03	0.48	
bathrooms	0.33	0.55	1.00	0.76	0.11	0.49	0.08	0.21	-0.12	0.69	
sqft_living	0.43	0.59	0.76	1.00	0.21	0.34	0.12	0.31	-0.06	0.88	
sqft_lot	0.05	0.07	0.11	0.21	1.00	0.00	0.02	0.07	0.00	0.22	
floors	0.15	0.18	0.49	0.34	0.00	1.00	0.02	0.03	-0.28	0.52	-
waterfront	0.14	-0.00	0.08	0.12	0.02	0.02	1.00	0.36	0.00	0.08	
view	0.23	0.11	0.21	0.31	0.07	0.03	0.36	1.00	0.06	0.17	
condition	0.03	0.03	-0.12	-0.06	0.00	-0.28	0.00	0.06	1.00	-0.18	
sqft_above	0.37	0.48	0.69	0.88	0.22	0.52	0.08	0.17	-0.18	1.00	-
sqft_basement	0.21	0.33	0.30	0.45	0.03	-0.26	0.10	0.32	0.20	-0.04	
yr_built	0.02	0.14	0.46	0.29	0.05	0.47	-0.02	-0.06	-0.40	0.41	-
yr_renovated	-0.03	-0.06	-0.22	-0.12	-0.02	-0.23	0.01	0.02	-0.19	-0.16	

There are correlation for price to all columns with positive correlation and negative with year renovated

Plot histogram


```
In [14]: # Plot histogram grid
df2.hist(figsize=(16,16), xrot=-45) ## Display the labels rotated by 45 degrees
S


# Clear the text "residue"
plt.show()
```



Using boxplot to see the outliers


```
In [15]: sns.boxplot(df2.sqft_lot)
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x15f6c348>
```




```
In [16]: sns.boxplot(df2.bedrooms)
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x16e7c688>
```




```
In [17]: sns.boxplot(df2.bathrooms)
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x16ed2848>
```



```
In [18]: sns.boxplot(df2.sqft_living)
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x167b54c8>
```



From the above it was seen that the sqft lot have so many outliers, so we need to do the cleaning by using filtering the data

Filtering the data and select the sqf lot below 500,000

```
In [19]: df2.sqft_lot.sort_values(ascending=False).head()
```

```
Out[19]: 1078      1074218
         2480      641203
         3487      478288
         375      435600
         879      435600
         Name: sqft_lot, dtype: int64
```


It was seen that there are range so high

```
In [20]: df2 = df2[df.sqft_lot <= 500000]
         df2.shape
```

```
Out[20]: (4598, 14)
```

```
In [21]: sns.boxplot(df2.sqft_lot)
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x16bd3948>
```



Now the outliers has reduced

One hot encoder for category data

```
In [22]: import category_encoders as ce
         # create an object of the OneHotEncoder
         OHE = ce.OneHotEncoder(cols=['city'],use_cat_names=True)
         # encode the categorical variables
         df3 = OHE.fit_transform(df2)
```

In [23]: `df3.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 4598 entries, 0 to 4599
```

```
Data columns (total 57 columns):
```

#	Column	Non-Null Count	Dtype
0	price	4598 non-null	float64
1	bedrooms	4598 non-null	float64
2	bathrooms	4598 non-null	float64
3	sqft_living	4598 non-null	int64
4	sqft_lot	4598 non-null	int64
5	floors	4598 non-null	float64
6	waterfront	4598 non-null	int64
7	view	4598 non-null	int64
8	condition	4598 non-null	int64
9	sqft_above	4598 non-null	int64
10	sqft_basement	4598 non-null	int64
11	yr_built	4598 non-null	int64
12	yr_renovated	4598 non-null	int64
13	city_Shoreline	4598 non-null	int64
14	city_Seattle	4598 non-null	int64
15	city_Kent	4598 non-null	int64
16	city_Bellevue	4598 non-null	int64
17	city_Redmond	4598 non-null	int64
18	city_Maple Valley	4598 non-null	int64
19	city_North Bend	4598 non-null	int64
20	city_Lake Forest Park	4598 non-null	int64
21	city_Sammamish	4598 non-null	int64
22	city_Auburn	4598 non-null	int64
23	city_Des Moines	4598 non-null	int64
24	city_Bothell	4598 non-null	int64
25	city_Federal Way	4598 non-null	int64
26	city_Kirkland	4598 non-null	int64
27	city_Issaquah	4598 non-null	int64
28	city_Woodinville	4598 non-null	int64
29	city_Normandy Park	4598 non-null	int64
30	city_Fall City	4598 non-null	int64
31	city_Renton	4598 non-null	int64
32	city_Carnation	4598 non-null	int64
33	city_Snoqualmie	4598 non-null	int64
34	city_Duvall	4598 non-null	int64
35	city_Burien	4598 non-null	int64
36	city_Covington	4598 non-null	int64
37	city_Inglewood-Finn Hill	4598 non-null	int64
38	city_Kenmore	4598 non-null	int64
39	city_Newcastle	4598 non-null	int64
40	city_Mercer Island	4598 non-null	int64
41	city_Black Diamond	4598 non-null	int64
42	city_Ravensdale	4598 non-null	int64
43	city_Clyde Hill	4598 non-null	int64
44	city_Algona	4598 non-null	int64
45	city_Skykomish	4598 non-null	int64
46	city_Tukwila	4598 non-null	int64
47	city_Vashon	4598 non-null	int64
48	city_Yarrow Point	4598 non-null	int64
49	city_SeaTac	4598 non-null	int64
50	city_Medina	4598 non-null	int64
51	city_Enumclaw	4598 non-null	int64

```

52 city_Snoqualmie Pass      4598 non-null    int64
53 city_Pacific              4598 non-null    int64
54 city_Beaux Arts Village   4598 non-null    int64
55 city_Preston              4598 non-null    int64
56 city_Milton               4598 non-null    int64
dtypes: float64(4), int64(53)
memory usage: 2.0 MB

```

Select the X and Y

In [24]: `df3.iloc[:, :12].describe()`

Out[24]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
count	4.598000e+03	4598.000000	4598.000000	4598.000000	4598.000000	4598.000000	4598.000000
mean	5.519002e+05	3.400826	2.160613	2139.127012	14485.896694	1.511962	0.000000
std	5.639402e+05	0.908505	0.783783	963.328580	30962.062220	0.538357	0.000000
min	0.000000e+00	0.000000	0.000000	370.000000	638.000000	1.000000	0.000000
25%	3.226250e+05	3.000000	1.750000	1460.000000	5000.250000	1.000000	0.000000
50%	4.604435e+05	3.000000	2.250000	1980.000000	7683.000000	1.500000	0.000000
75%	6.547125e+05	4.000000	2.500000	2620.000000	11000.000000	2.000000	0.000000
max	2.659000e+07	9.000000	8.000000	13540.000000	478288.000000	3.500000	1.000000

Need to do scaling the data using data scaler due to different in standard deviation and mean

In [25]: `y = df3['price']`

In [26]: `X=df3.iloc[:, 1:57]`

In [27]: `X.shape`

Out[27]: (4598, 56)

In [28]: `y.shape`

Out[28]: (4598,)

Use the datascaler

```

In [29]: from sklearn.preprocessing import StandardScaler
scaler_x = StandardScaler()
X_scaled = scaler_x.fit_transform(X)

```

```
In [30]: print(X_scaled)
```

```
[[-0.44124149 -0.84294362 -0.82963793 ... -0.01474901 -0.02086051
  -0.02086051]
 [ 1.76041708  0.43305795  1.56855858 ... -0.01474901 -0.02086051
  -0.02086051]
 [-0.44124149 -0.20494284 -0.21711155 ... -0.01474901 -0.02086051
  -0.02086051]
 ...
 [-0.44124149  0.43305795  0.90412319 ... -0.01474901 -0.02086051
  -0.02086051]
 [ 0.6595878  -0.20494284 -0.0510027  ... -0.01474901 -0.02086051
  -0.02086051]
 [-0.44124149  0.43305795 -0.67391088 ... -0.01474901 -0.02086051
  -0.02086051]]
```

```
In [31]: X_scaled.shape
```

```
Out[31]: (4598, 56)
```

```
In [32]: y.shape
```

```
Out[32]: (4598,)
```

```
In [33]: y = y.values.reshape(-1,1)
```

```
In [34]: scaler_y = StandardScaler()
y_scaled = scaler_y.fit_transform(y)
```

Split the dataset

```
In [35]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_scaled, test_s
size = .25, random_state = 3)
```

Regression Model

1. Linear Regression as baseline

```
In [36]: from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

```
Out[36]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [37]: y_pred_lin = regressor.predict(X_test)
```



```
In [38]: y_train_lin = regressor.predict(X_train)
```

```
In [39]: y_predict_lin = scaler_y.inverse_transform(y_pred_lin)
```

```
In [40]: y_train_linear = scaler_y.inverse_transform(y_train_lin)
```

```
In [41]: y_test_orig = scaler_y.inverse_transform(y_test)
```

```
In [42]: y_train_orig = scaler_y.inverse_transform(y_train)
```

```
In [43]: print("Train Results for Linear regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_train_orig, y_train_linear)))
print("R-squared: ", rs(y_train_orig, y_train_linear))
print("Mean Absolute Error: ", mae(y_train_orig, y_train_linear))
```

```
Train Results for Linear regression:
*****
Root mean squared error:  543776.6926778575
R-squared:  0.23306523863088413
Mean Absolute Error:  143356.0830687553
```

```
In [92]: RMSE_lin = sqrt(mse(y_train_orig, y_train_linear))
R_square_lin = rs(y_train_orig, y_train_linear)
MAE_lin = mae(y_train_orig, y_train_linear)
```

```
In [44]: print("Test Results for Linear regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_test_orig, y_predict_lin)))
print("R-squared: ", rs(y_test_orig, y_predict_lin))
print("Mean Absolute Error: ", mae(y_test_orig, y_predict_lin))
```

```
Test Results for Linear regression:
*****
Root mean squared error:  217349.7619004028
R-squared:  0.5892605775944486
Mean Absolute Error:  129180.659933383
```

Ridge Regression

```
In [45]: tuned_params = {'alpha': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000, 100000]}
model_Ridge = GridSearchCV(Ridge(), tuned_params, scoring = 'neg_mean_absolute_error', cv=10, n_jobs=-1)
model_Ridge.fit(X_train, y_train)
```

```
Out[45]: GridSearchCV(cv=10, error_score=nan,
                      estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
                                      max_iter=None, normalize=False, random_state=None,
                                      solver='auto', tol=0.001),
                      iid='deprecated', n_jobs=-1,
                      param_grid={'alpha': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000, 100000]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring='neg_mean_absolute_error', verbose=0)
```

```
In [46]: model_Ridge.best_estimator_
```

```
Out[46]: Ridge(alpha=1000, copy_X=True, fit_intercept=True, max_iter=None,
               normalize=False, random_state=None, solver='auto', tol=0.001)
```

```
In [47]: ## Predict Train results
y_train_ridge = model_Ridge.predict(X_train)
```

```
In [48]: ## Predict Test results
y_pred_ridge = model_Ridge.predict(X_test)
```

```
In [49]: y_predict_ridge = scaler_y.inverse_transform(y_pred_ridge)
```

```
In [50]: y_train_ridge = scaler_y.inverse_transform(y_train_ridge)
```

```
In [51]: print("Train Results for Ridge regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_train_orig, y_train_ridge)))
print("R-squared: ", rs(y_train_orig, y_train_ridge))
print("Mean Absolute Error: ", mae(y_train_orig, y_train_ridge))
```

```
Train Results for Ridge regression:
*****
Root mean squared error:  545922.5975553852
R-squared:  0.22700018871270244
Mean Absolute Error:  140999.8893371432
```

```
In [93]: RMSE_ridge = sqrt(mse(y_train_orig, y_train_ridge))
R_square_ridge = rs(y_train_orig, y_train_ridge)
MAE_ridge = mae(y_train_orig, y_train_ridge)
```

```
In [52]: print("Test Results for Ridge regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_test_orig, y_predict_ridge)))
print("R-squared: ", rs(y_test_orig, y_predict_ridge))
print("Mean Absolute Error: ", mae(y_test_orig, y_predict_ridge))
```

Test Results for Ridge regression:

Root mean squared error: 216476.82764744837

R-squared: 0.5925532289854684

Mean Absolute Error: 126089.03185500484

Support Vector

```
In [53]: ## Building the model again with the best hyperparameters
model_SVR = SVR(kernel='linear')
model_SVR.fit(X_train, y_train)
```

C:\Users\BIrawan\Anaconda3\lib\site-packages\sklearn\utils\validation.py:760:
DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

```
Out[53]: SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1, gamma='scale',
kernel='linear', max_iter=-1, shrinking=True, tol=0.001, verbose=False)
```

```
In [54]: ## Predict Train results
y_train_SVR = model_SVR.predict(X_train)
```

```
In [55]: ## Predict Test results
y_pred_SVR = model_SVR.predict(X_test)
```

```
In [56]: y_predict_SVR = scaler_y.inverse_transform(y_pred_SVR)
```

```
In [57]: y_trainSVR = scaler_y.inverse_transform(y_train_SVR)
```

```
In [58]: print("Train Results for Support Vector regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_train_orig, y_trainSVR)))
print("R-squared: ", rs(y_train_orig, y_trainSVR))
print("Mean Absolute Error: ", mae(y_train_orig, y_trainSVR))
```

Train Results for Support Vector regression:

Root mean squared error: 548026.8147132694

R-squared: 0.22102976653869932

Mean Absolute Error: 131408.60806086232

```
In [94]: RMSE_svr = sqrt(mse(y_train_orig, y_trainSVR))
R_square_svr = rs(y_train_orig, y_trainSVR)
MAE_svr = mae(y_train_orig, y_trainSVR)
```

```
In [59]: print("Test Results for Support Vector regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_test_orig, y_predict_SVR)))
print("R-squared: ", rs(y_test_orig, y_predict_SVR))
print("Mean Absolute Error: ", mae(y_test_orig, y_predict_SVR))
```

Test Results for Support Vector regression:

Root mean squared error: 214217.2718139181

R-squared: 0.6010145865804548

Mean Absolute Error: 116611.49934353186

Random Forest

```
In [60]: tuned_params = {'n_estimators': [100, 200, 300, 400, 500], 'min_samples_split'
: [2, 5, 10], 'min_samples_leaf': [1, 2, 4]}
model_RF = RandomizedSearchCV(RandomForestRegressor(), tuned_params, n_iter=20
, scoring = 'neg_mean_absolute_error', cv=5, n_jobs=-1)
model_RF.fit(X_train, y_train)
```

C:\Users\BIrawan\Anaconda3\lib\site-packages\sklearn\model_selection_search.py:739: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
self.best_estimator_.fit(X, y, **fit_params)
```

```
Out[60]: RandomizedSearchCV(cv=5, error_score=nan,
        estimator=RandomForestRegressor(bootstrap=True,
                                         ccp_alpha=0.0,
                                         criterion='mse',
                                         max_depth=None,
                                         max_features='auto',
                                         max_leaf_nodes=None,
                                         max_samples=None,
                                         min_impurity_decrease=0.0,
                                         min_impurity_split=None,
                                         min_samples_leaf=1,
                                         min_samples_split=2,
                                         min_weight_fraction_leaf=
0.0,
                                         n_estimators=100,
                                         n_jobs=None, oob_score=False,
                                         random_state=None, verbose=
0,
                                         warm_start=False),
        iid='deprecated', n_iter=20, n_jobs=-1,
        param_distributions={'min_samples_leaf': [1, 2, 4],
                             'min_samples_split': [2, 5, 10],
                             'n_estimators': [100, 200, 300, 400,
                                                500]},
        pre_dispatch='2*n_jobs', random_state=None, refit=True,
        return_train_score=False, scoring='neg_mean_absolute_error',
        verbose=0)
```

```
In [61]: model_RF.best_estimator_
```

```
Out[61]: RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                max_samples=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=10, min_weight_fraction_leaf=0.0,
                                n_estimators=100, n_jobs=None, oob_score=False,
                                random_state=None, verbose=0, warm_start=False)
```

```
In [62]: ## Predict Train results
y_train_RF = model_RF.predict(X_train)
```

```
In [63]: ## Predict Test results
y_pred_RF = model_RF.predict(X_test)
```

```
In [64]: y_predict_RF= scaler_y.inverse_transform(y_pred_RF)
```

```
In [65]: y_trainRF = scaler_y.inverse_transform(y_train_RF)
```

```
In [66]: print("Train Results for Random Forrest regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_train_orig, y_trainRF)))
print("R-squared: ", rs(y_train_orig, y_trainRF))
print("Mean Absolute Error: ", mae(y_train_orig, y_trainRF))
```

```
Train Results for Random Forrest regression:
*****
Root mean squared error: 359409.76897367387
R-squared: 0.6649597381787381
Mean Absolute Error: 83788.20857374588
```

```
In [95]: RMSE_rf = sqrt(mse(y_train_orig, y_trainRF))
R_square_rf = rs(y_train_orig, y_trainRF)
MAE_rf = mae(y_train_orig, y_trainRF)
```

```
In [67]: print("Test Results for Random Forrest regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_test_orig, y_predict_RF)))
print("R-squared: ", rs(y_test_orig, y_predict_RF))
print("Mean Absolute Error: ", mae(y_test_orig, y_predict_RF))
```

```
Test Results for Random Forrest regression:
*****
Root mean squared error: 301957.37700281135
R-squared: 0.20724436301074567
Mean Absolute Error: 130711.47366073588
```

Decission Tree

```
In [68]: tuned_params = {'min_samples_split': [2, 3, 4, 5, 7], 'min_samples_leaf': [1,
2, 3, 4, 6], 'max_depth': [2, 3, 4, 5, 6, 7]}
model_DT = RandomizedSearchCV(DecisionTreeRegressor(), tuned_params, n_iter=20
, scoring = 'neg_mean_absolute_error', cv=10, n_jobs=-1)
model_DT.fit(X_train, y_train)
```

```
Out[68]: RandomizedSearchCV(cv=10, error_score=nan,
                             estimator=DecisionTreeRegressor(ccp_alpha=0.0,
                                                                criterion='mse',
                                                                max_depth=None,
                                                                max_features=None,
                                                                max_leaf_nodes=None,
                                                                min_impurity_decrease=0.0,
                                                                min_impurity_split=None,
                                                                min_samples_leaf=1,
                                                                min_samples_split=2,
                                                                min_weight_fraction_leaf=
0.0,
                                                                presort='deprecated',
                                                                random_state=None,
                                                                splitter='best'),
                             iid='deprecated', n_iter=20, n_jobs=-1,
                             param_distributions={'max_depth': [2, 3, 4, 5, 6, 7],
                                                  'min_samples_leaf': [1, 2, 3, 4, 6],
                                                  'min_samples_split': [2, 3, 4, 5,
7]},
                             pre_dispatch='2*n_jobs', random_state=None, refit=True,
                             return_train_score=False, scoring='neg_mean_absolute_erro
r',
                             verbose=0)
```

```
In [69]: model_DT.best_estimator_
```

```
Out[69]: DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=7,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=6, min_samples_split=5,
                                min_weight_fraction_leaf=0.0, presort='deprecated',
                                random_state=None, splitter='best')
```

```
In [70]: ## Predict Train results
y_train_DT = model_DT.predict(X_train)
```

```
In [71]: ## Predict Test results
y_pred_DT = model_DT.predict(X_test)
```

```
In [72]: y_predict_DT= scaler_y.inverse_transform(y_pred_DT)
```

```
In [73]: y_trainDT = scaler_y.inverse_transform(y_train_DT)
```

```
In [74]: print("Train Results for Random Forrest regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_train_orig, y_trainDT)))
print("R-squared: ", rs(y_train_orig, y_trainDT))
print("Mean Absolute Error: ", mae(y_train_orig, y_trainDT))
```

```
Train Results for Random Forrest regression:
*****
Root mean squared error:  503441.21157148166
R-squared:  0.3426225951880927
Mean Absolute Error:  155599.49303935387
```

```
In [96]: RMSE_dt = sqrt(mse(y_train_orig, y_trainDT))
R_square_dt = rs(y_train_orig, y_trainDT)
MAE_dt = mae(y_train_orig, y_trainDT)
```

```
In [75]: print("Test Results for Support Vector regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_test_orig, y_predict_DT)))
print("R-squared: ", rs(y_test_orig, y_predict_DT))
print("Mean Absolute Error: ", mae(y_test_orig, y_predict_DT))
```

```
Test Results for Support Vector regression:
*****
Root mean squared error:  367929.44616630883
R-squared:  -0.17700184590625723
Mean Absolute Error:  169146.10620063593
```

KNN

```
In [76]: # creating odd list of K for KNN
neighbors = list(range(1,50,2))
# empty list that will hold cv scores
cv_scores = []

# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsRegressor(n_neighbors=k)
    scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='neg_mean_absolute_error')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)
```

The optimal number of neighbors is 5.


```
In [77]: model_KNN = KNeighborsRegressor(n_neighbors = optimal_k)
        model_KNN.fit(X_train, y_train)
```

```
Out[77]: KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                             weights='uniform')
```

```
In [78]: ## Predict Train results
        y_train_KNN = model_KNN.predict(X_train)
```

```
In [79]: ## Predict Test results
        y_pred_KNN = model_KNN.predict(X_test)
```

```
In [80]: y_predict_KNN= scaler_y.inverse_transform(y_pred_KNN)
```

```
In [81]: y_trainKNN = scaler_y.inverse_transform(y_train_KNN)
```

```
In [82]: print("Train Results for KNN regression:")
        print("*****")
        print("Root mean squared error: ", sqrt(mse(y_train_orig, y_trainKNN)))
        print("R-squared: ", rs(y_train_orig, y_trainKNN))
        print("Mean Absolute Error: ", mae(y_train_orig, y_trainKNN))
```

```
Train Results for KNN regression:
*****
Root mean squared error:  475448.93609199184
R-squared:  0.41369310089289824
Mean Absolute Error:  120781.00356891201
```

```
In [97]: RMSE_knn = sqrt(mse(y_train_orig, y_trainKNN))
        R_square_knn = rs(y_train_orig, y_trainKNN)
        MAE_knn = mae(y_train_orig, y_trainKNN)
```

```
In [83]: print("Test Results for KNN regression:")
        print("*****")
        print("Root mean squared error: ", sqrt(mse(y_test_orig, y_predict_KNN)))
        print("R-squared: ", rs(y_test_orig, y_predict_KNN))
        print("Mean Absolute Error: ", mae(y_test_orig, y_predict_KNN))
```

```
Test Results for KNN regression:
*****
Root mean squared error:  293743.8430326016
R-squared:  0.24978525557596376
Mean Absolute Error:  133455.36335048106
```

Lasso Regression

```
In [84]: tuned_params = {'alpha': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000, 100000]}
model_LS = GridSearchCV(Lasso(), tuned_params, scoring = 'neg_mean_absolute_error', cv=20, n_jobs=-1)
model_LS.fit(X_train, y_train)
```

```
Out[84]: GridSearchCV(cv=20, error_score=nan,
                    estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True,
                                    max_iter=1000, normalize=False, positive=False,
                                    precompute=False, random_state=None,
                                    selection='cyclic', tol=0.0001, warm_start=False),
                    iid='deprecated', n_jobs=-1,
                    param_grid={'alpha': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000,
                                           10000, 100000]}},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                    scoring='neg_mean_absolute_error', verbose=0)
```

```
In [85]: model_LS.best_estimator_
```

```
Out[85]: Lasso(alpha=0.01, copy_X=True, fit_intercept=True, max_iter=1000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
```

```
In [86]: ## Predict Train results
y_train_LS = model_LS.predict(X_train)
```

```
In [87]: ## Predict Test results
y_pred_LS = model_LS.predict(X_test)
```

```
In [88]: y_predict_LS= scaler_y.inverse_transform(y_pred_LS)
```

```
In [89]: y_trainLS = scaler_y.inverse_transform(y_train_LS)
```

```
In [98]: print("Train Results for Lasso regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_train_orig, y_trainLS)))
print("R-squared: ", rs(y_train_orig, y_trainLS))
print("Mean Absolute Error: ", mae(y_train_orig, y_trainLS))
```

```
Train Results for Lasso regression:
*****
Root mean squared error:  544980.8625046085
R-squared:  0.22966479040177779
Mean Absolute Error:  142639.46170476923
```

```
In [100]: RMSE_ls = sqrt(mse(y_train_orig, y_trainLS))
R_square_ls = rs(y_train_orig, y_trainLS)
MAE_ls = mae(y_train_orig, y_trainLS)
```

```
In [99]: print("Test Results for Lasso regression:")
print("*****")
print("Root mean squared error: ", sqrt(mse(y_test_orig, y_predict_LS)))
print("R-squared: ", rs(y_test_orig, y_predict_LS))
print("Mean Absolute Error: ", mae(y_test_orig, y_predict_LS))
```

Test Results for Lasso regression:

Root mean squared error: 217825.14016566932

R-squared: 0.5874619087493652

Mean Absolute Error: 127931.71645607855

CREATE TABLE OF COMPARISSON FOR ALL MODELS

```
In [108]: data_comp = {'Model Name': ['Linear Regression', 'Ridge Regression', 'Support
Vector', 'Random Forrest', 'KNN', 'Lasso'],
                      'RMSE': [RMSE_lin, RMSE_ridge, RMSE_svr, RMSE_rf, RMSE_knn, RMSE_ls],
                      'R-Squared': [R_square_lin, R_square_ridge, R_square_svr, R_square_rf, R_s
quare_knn, R_square_ls],
                      'MAE': [MAE_lin, MAE_ridge, MAE_svr, MAE_rf, MAE_knn, MAE_ls]
                      }

df_comp = pd.DataFrame (data_comp, columns = ['Model Name', 'RMSE', 'R-Squared'
, 'MAE'])

print (df_comp)
```

	Model Name	RMSE	R-Squared	MAE
0	Linear Regression	543776.692678	0.233065	143356.083069
1	Ridge Regression	545922.597555	0.227000	140999.889337
2	Support Vector	548026.814713	0.221030	131408.608061
3	Random Forrest	359409.768974	0.664960	83788.208574
4	KNN	475448.936092	0.413693	120781.003569
5	Lasso	544980.862505	0.229665	142639.461705

```
In [110]: df_comp.sort_values(by='R-Squared', ascending=False )
```

Out[110]:

	Model Name	RMSE	R-Squared	MAE
3	Random Forrest	359409.768974	0.664960	83788.208574
4	KNN	475448.936092	0.413693	120781.003569
0	Linear Regression	543776.692678	0.233065	143356.083069
5	Lasso	544980.862505	0.229665	142639.461705
1	Ridge Regression	545922.597555	0.227000	140999.889337
2	Support Vector	548026.814713	0.221030	131408.608061

CONCLUSION

From the above table it was shown that the Random Forrest is the best regression model on predicting the housing price in Boston dataset, due to :

1. highest score in R-squared
2. Lowest score in RMSE and MAE

Further improvement :

1. Using L1 and L2

Furthermore, the pipeline model can be made using random forest